

A scatter plot showing the training loss over 3500 episodes. The x-axis is labeled 'trained episodes' and ranges from 0 to 3500 with major ticks every 500 units. The y-axis represents the loss, with a horizontal line at 0.0. The data points are blue dots. The plot shows a noisy training process with a general downward trend. The loss starts around 0.15 at episode 0, fluctuates between 0.05 and 0.15 until episode 1000, and then gradually decreases to around 0.05 by episode 2000. After episode 2000, the loss stabilizes around 0.05, with a few outliers reaching up to 0.1. The overall trend suggests that the model is learning and improving over time, despite the high variance in the loss values.

